

# Lab: Accelerating Code: From Python Baseline to Architecture-Aware Optimization

## Dense Matrix-Matrix Multiplication (GEMM) Challenge

Team Name: [Your Team Name Here]  
Entry Numbers: [mcsXXXXX, mcsYYYYY]

November 15, 2025

### Abstract

This report details the iterative process of optimizing a Dense Matrix-Matrix Multiplication (GEMM) algorithm, starting from a Python baseline. We describe the implementation of several architecture-aware techniques, including OpenMP for multicore parallelism, cache-friendly blocking, and SIMD vectorization using AVX intrinsics. We present a comprehensive performance analysis, measuring GFLOPS, speedup, and scalability against thread count and problem size. Our final optimized C++ implementation achieves a significant speedup of [X.XX]x over the baseline, demonstrating the performance impact of architecture-level design.

## Contents

<b>1 Problem Description</b>	<b>2</b>
<b>2 Baseline Implementation</b>	<b>2</b>
<b>3 Optimizations Implemented</b>	<b>2</b>
3.1 Multithreading with OpenMP . . . . .	2
3.2 Cache-Friendly Blocking . . . . .	2
3.3 SIMD Vectorization . . . . .	3
3.4 Other Optimizations . . . . .	3
<b>4 Experimental Methodology</b>	<b>3</b>
<b>5 Performance Results</b>	<b>3</b>

# 1 Problem Description

The objective of this assignment is to accelerate a Dense Matrix-Matrix Multiplication (GEMM) operation,  $C = A \times B$ . The baseline implementation is provided in Python, using `numpy` and the `multiprocessing` library. Our goal is to reimplement this in C++ and apply a series of architecture-aware optimizations to maximize performance, measured in GFLOPS and speedup.

This report documents our optimizations for input matrix sizes

$$N \in \{512, 1024, 1536, 2048\}.$$

## 2 Baseline Implementation

The baseline consists of the `gemm_baseline.py` script. It uses the Python `multiprocessing` library to distribute blocks of matrix A to worker processes, where each worker computes its portion using `np.dot()`. Although `numpy` is optimized, the Python overhead is significant.

Our environment:

- **CPU Model:** [Your CPU Model]
- **Cores/Threads:** [e.g., 16 cores / 32 threads]
- **Caches:** [e.g., L1/L2/L3 = 32K/1024K/36608K]
- **OS:** [e.g., Ubuntu 20.04.5 LTS]
- **Compiler:** g++ 11.3.0

## 3 Optimizations Implemented

We implemented the following optimizations.

### 3.1 Multithreading with OpenMP

A major bug was fixed—OpenMP was incorrectly placed inside inner loops. The final correct placement:


```
// Correct OpenMP structure
#pragma omp parallel for schedule(static)
for (int i0 = 0; i0 < N; i0 += MC) {
    for (int k0 = 0; k0 < N; k0 += KC) {
        for (int j0 = 0; j0 < N; j0 += NC) {
            // ...
        }
    }
}
```

### 3.2 Cache-Friendly Blocking

To reduce cache misses, we divided matrices into blocks:

- A:  $MC \times KC$  - B:  $KC \times NC$  - C:  $MC \times NC$

We chose:  $MC = 128$ ,  $KC = 256$ ,  $NC = 256$ .



**Image Placeholder: Cache Blocking Diagram**

*Block-level diagram showing  $i0$ ,  $k0$ ,  $j0$  tiling.*

Figure 1: Our cache-blocking strategy.

### 3.3 SIMD Vectorization

We implemented an AVX2/AVX-512 microkernel to compute the innermost block fully in vector registers.

### 3.4 Other Optimizations

- Pre-transposition of matrix B for sequential memory access.
- 64-byte aligned memory for SIMD.
- Compiler flags: `-O3 -march=native -fopenmp`.

## 4 Experimental Methodology

We used automated scripts:

1. Warmup run
2. 10 measured runs per configuration
3. Logged to `gemm_results.csv`

## 5 Performance Results

Table 1: Key Performance Metrics (Avg. of 10 Runs)

Metric	512	1024	1536	2048
Baseline Time (s)	[ ]	[ ]	[ ]	[ ]
Optimized Time (s, T=16)	[ ]	[ ]	[ ]	[ ]
Optimized GFLOPS	[ ]	[ ]	[ ]	[ ]
Max Speedup	[ ]	[ ]	[ ]	[ ]

**Image Placeholder: GFLOPS vs Threads Plot**

Figure 2: Performance (GFLOPS) vs. Thread Count.



Figure 3: Speedup vs. Thread Count.